

# Bushfire Hazard Reduction Feasibility Prediction

Randompdf

Exported and Printed

February 6, 2025

## 0.1 Introduction

This document provides an **overview of the Bushfire Hazard Reduction Feasibility Prediction project**, which applies **machine learning techniques** to predict whether a controlled burn can be conducted safely based on historical weather data. The objective is to support decision-making in mitigating bushfire risks in Australia.

The study is based on a dataset containing **over 134,000 records with 24 variables**, covering **weather conditions, location, and seasonal factors**. A structured approach was followed, including **data preprocessing, feature engineering, class imbalance handling, and model evaluation**.

### 0.1.1 Exploratory Data Analysis and Clustering

As part of the initial analysis, **k-Means Clustering** was applied to identify potential patterns in the data. However, the **Silhouette Score of 0.213** indicated that clustering did not provide strong separation for classification purposes.

### 0.1.2 Classification Algorithms Used

Several machine learning models were employed for classification, including:

- **Decision Tree**
- **Artificial Neural Network (ANN)**
- **Gaussian Naïve Bayes (NBG)**
- **Support Vector Machine (SVM)**
- **Random Forest**

Each model was evaluated using **accuracy, precision, recall, F1-score, and confusion matrices**, with a special focus on **handling class imbalance** and optimizing performance. Among the models, **Random Forest demonstrated the highest accuracy and balanced precision-recall, making it a strong candidate for practical deployment**.

### 0.1.3 Access to Full Project

For those interested in further exploration, the **full code, environment setup, and dataset** are available **upon request**. If you would like access to these materials, please reach out.

This document will also be uploaded to a **public GitHub repository**, ensuring transparency and accessibility for the research and machine learning communities.

```
[1]: # The order of importing libraries are according to the sequence of steps taken
      ↪ in our work.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import openpyxl

from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer

from sklearn.model_selection import train_test_split

from sklearn.neural_network import MLPClassifier

from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import classification_report,
      ↪ confusion_matrix, accuracy_score
from sklearn.model_selection import cross_val_score
```

```
[2]: # Load dataset
df = pd.read_excel("A1_Data.xlsx")
```

```
[3]: df.shape
```

```
[3]: (134692, 24)
```

```
[4]: # Get the data types of each column
      column_types = df.dtypes

      # Print the data types
      print(column_types)
```

|                     |         |
|---------------------|---------|
| City_Name           | object  |
| City_State          | object  |
| Date                | object  |
| Season              | object  |
| Climate             | object  |
| Minimum_Temperature | float64 |
| Maximum_Temperature | float64 |
| Temperature_AM      | float64 |
| Temperature_PM      | float64 |
| Wind_Direction_AM   | object  |
| Wind_Direction_PM   | object  |
| Wind_Speed_AM       | float64 |
| Wind_Speed_PM       | float64 |

```

Max_WindGust_Direction      object
Max_WindGust_Speed          float64
Significant_Rainfall        object
Rainfall_Amount             float64
Humidity_AM                 float64
Humidity_PM                 float64
Atmospheric_Pressure_AM     float64
Atmospheric_Pressure_PM     float64
CloudCover_Oktas_AM         int64
CloudCover_Oktas_PM         int64
Burn_Tomorrow               int64
dtype: object

```

```

[5]: def identify_data(df):
      # Check for missing values
      print('Missing values in the DataFrame:\n', df.isnull().sum(),'\n')
      # Check for duplicates
      print('Number of duplicated rows in the DataFrame:\n', df.duplicated().
      ↪sum(),'\n')
      # Check for unique values
      print('Number of unique values in each column of the DataFrame:\n', df.
      ↪nunique(),'\n')
      # Check for out of bound value
      stats=df.describe()
      print(stats)

      # Set the figure size to make the plot wider
      plt.figure(figsize=(16, 6))
      #Create a boxplot for each numeric column
      ax = df.boxplot()

      # Set the rotation angle for x-axis labels
      plt.xticks(rotation=45) # You can adjust the rotation angle as needed

      # Show the plot
      plt.show()
      identify_data(df)

```

Missing values in the DataFrame:

```

City_Name      0
City_State      0
Date           0
Season         0
Climate        0
Minimum_Temperature  582
Maximum_Temperature  305
Temperature_AM  0

```

|                         |      |
|-------------------------|------|
| Temperature_PM          | 0    |
| Wind_Direction_AM       | 9531 |
| Wind_Direction_PM       | 3558 |
| Wind_Speed_AM           | 1323 |
| Wind_Speed_PM           | 2452 |
| Max_WindGust_Direction  | 9027 |
| Max_WindGust_Speed      | 8968 |
| Significant_Rainfall    | 0    |
| Rainfall_Amount         | 1350 |
| Humidity_AM             | 0    |
| Humidity_PM             | 0    |
| Atmospheric_Pressure_AM | 0    |
| Atmospheric_Pressure_PM | 0    |
| CloudCover_Oktas_AM     | 0    |
| CloudCover_Oktas_PM     | 0    |
| Burn_Tomorrow           | 0    |

dtype: int64

Number of duplicated rows in the DataFrame:  
0

Number of unique values in each column of the DataFrame:

|                         |      |
|-------------------------|------|
| City_Name               | 48   |
| City_State              | 8    |
| Date                    | 3337 |
| Season                  | 4    |
| Climate                 | 5    |
| Minimum_Temperature     | 389  |
| Maximum_Temperature     | 505  |
| Temperature_AM          | 440  |
| Temperature_PM          | 500  |
| Wind_Direction_AM       | 16   |
| Wind_Direction_PM       | 16   |
| Wind_Speed_AM           | 43   |
| Wind_Speed_PM           | 44   |
| Max_WindGust_Direction  | 16   |
| Max_WindGust_Speed      | 67   |
| Significant_Rainfall    | 2    |
| Rainfall_Amount         | 666  |
| Humidity_AM             | 165  |
| Humidity_PM             | 166  |
| Atmospheric_Pressure_AM | 545  |
| Atmospheric_Pressure_PM | 548  |
| CloudCover_Oktas_AM     | 10   |
| CloudCover_Oktas_PM     | 10   |
| Burn_Tomorrow           | 2    |

dtype: int64

|       | Minimum_Temperature | Maximum_Temperature | Temperature_AM \ |
|-------|---------------------|---------------------|------------------|
| count | 134110.000000       | 134387.000000       | 134692.000000    |
| mean  | 12.144484           | 23.321705           | 16.970687        |
| std   | 6.428055            | 7.214363            | 6.554124         |
| min   | -8.500000           | -4.800000           | -7.200000        |
| 25%   | 7.500000            | 17.900000           | 12.200000        |
| 50%   | 11.900000           | 22.800000           | 16.600000        |
| 75%   | 16.800000           | 28.500000           | 21.600000        |
| max   | 33.900000           | 48.100000           | 40.200000        |


|       | Temperature_PM | Wind_Speed_AM | Wind_Speed_PM | Max_WindGust_Speed \ |
|-------|----------------|---------------|---------------|----------------------|
| count | 134692.000000  | 133369.000000 | 132240.000000 | 125724.000000        |
| mean  | 21.792406      | 13.951293     | 18.663203     | 40.126173            |
| std   | 7.037912       | 8.892540      | 8.836756      | 13.598228            |
| min   | -5.400000      | 0.000000      | 0.000000      | 6.000000             |
| 25%   | 16.600000      | 7.000000      | 13.000000     | 31.000000            |
| 50%   | 21.200000      | 13.000000     | 19.000000     | 39.000000            |
| 75%   | 26.700000      | 19.000000     | 24.000000     | 48.000000            |
| max   | 46.700000      | 130.000000    | 87.000000     | 135.000000           |

|       | Rainfall_Amount | Humidity_AM   | Humidity_PM   | Atmospheric_Pressure_AM \ |
|-------|-----------------|---------------|---------------|---------------------------|
| count | 133342.000000   | 134692.000000 | 134692.000000 | 134692.000000             |
| mean  | 2.344291        | 68.645683     | 51.139806     | 1017.512641               |
| std   | 8.472936        | 19.079176     | 20.719840     | 6.775998                  |
| min   | 0.000000        | 0.000000      | 0.000000      | 982.000000                |
| 25%   | 0.000000        | 57.000000     | 36.000000     | 1013.300000               |
| 50%   | 0.000000        | 70.000000     | 52.000000     | 1017.900000               |
| 75%   | 0.600000        | 83.000000     | 65.000000     | 1021.600000               |
| max   | 371.000000      | 100.000000    | 100.000000    | 1041.000000               |

|       | Atmospheric_Pressure_PM | CloudCover_Oktas_AM | CloudCover_Oktas_PM \ |
|-------|-------------------------|---------------------|-----------------------|
| count | 134692.000000           | 134692.000000       | 134692.000000         |
| mean  | 1015.108536             | 4.517492            | 4.523015              |
| std   | 6.709368                | 2.375169            | 2.189380              |
| min   | 977.100000              | 0.000000            | 0.000000              |
| 25%   | 1010.800000             | 3.000000            | 3.000000              |
| 50%   | 1015.600000             | 4.000000            | 4.000000              |
| 75%   | 1019.200000             | 6.000000            | 6.000000              |
| max   | 1039.600000             | 9.000000            | 9.000000              |

|       | Burn_Tomorrow |
|-------|---------------|
| count | 134692.000000 |
| mean  | 0.277648      |
| std   | 0.447841      |
| min   | 0.000000      |
| 25%   | 0.000000      |
| 50%   | 0.000000      |
| 75%   | 1.000000      |

max 1.000000



```
[6]: # Extract the month from the 'Date' column
df['Date'] = pd.to_datetime(df['Date'])
df['Month'] = df['Date'].dt.month

# Columns for imputation
numerical_columns_to_impute = ['Minimum_Temperature', 'Maximum_Temperature', 'Wind_Speed_AM', 'Wind_Speed_PM', 'Rainfall_Amount', 'Max_WindGust_Speed']
categorical_columns_to_impute = ['Wind_Direction_AM', 'Wind_Direction_PM', 'Max_WindGust_Direction', 'Climate', 'Season', 'Significant_Rainfall']

# Impute numerical columns with monthly median
median_imputer = SimpleImputer(strategy='median')
for col in numerical_columns_to_impute:
    df[col] = df.groupby('Month')[col].transform(lambda x: median_imputer.fit_transform(x.values.reshape(-1, 1)).flatten())

# Impute categorical columns with monthly mode
mode_imputer = SimpleImputer(strategy='most_frequent')
for col in categorical_columns_to_impute:
    df[col] = df.groupby('Month')[col].transform(lambda x: mode_imputer.fit_transform(x.values.reshape(-1, 1)).flatten())

# Drop the 'Month' column not needed anymore
df.drop('Month', axis=1, inplace=True)
```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[6], line 2
      1 # Extract the month from the 'Date' column
----> 2 df['Date'] = pd.to_datetime(df['Date'])
      3 df['Month'] = df['Date'].dt.month
      5 # Columns for imputation

File E:\Python\Lib\site-packages\pandas\core\tools\datetimes.py:1112, in
    to_datetime(arg, errors, dayfirst, yearfirst, utc, format, exact, unit,
    infer_datetime_format, origin, cache)
    1110         result = arg.map(cache_array)
    1111     else:
-> 1112         values = convert_listlike(arg._values, format)
    1113         result = arg._constructor(values, index=arg.index, name=arg.name)
    1114 elif isinstance(arg, (ABCDDataFrame, abc.MutableMapping)):

File E:\Python\Lib\site-packages\pandas\core\tools\datetimes.py:488, in
    _convert_listlike_datetimes(arg, format, name, utc, unit, errors, dayfirst,
    yearfirst, exact)
    486 # `format` could be inferred, or user didn't ask for mixed-format
    parsing.
    487 if format is not None and format != "mixed":
--> 488     return
    _array_strptime_with_fallback(arg, name, utc, format, exact, errors)
    490 result, tz_parsed = objects_to_datetime64ns(
    491     arg,
    492     dayfirst=dayfirst,
    (...)
    496     allow_object=True,
    497 )
    499 if tz_parsed is not None:
    500     # We can take a shortcut since the datetime64 numpy array
    501     # is in UTC

File E:\Python\Lib\site-packages\pandas\core\tools\datetimes.py:519, in
    _array_strptime_with_fallback(arg, name, utc, fmt, exact, errors)
    508 def _array_strptime_with_fallback(
    509     arg,
    510     name,
    (...)
    514     errors: str,
    515 ) -> Index:
    516     """
    517     Call array_strptime, with fallback behavior depending on 'errors'.
    518     """

```

```

--> 519     result, timezones = _
    ↪ array_strptime(arg, fmt, exact=exact, errors=errors, utc=utc)
    520     if any(tz is not None for tz in timezones):
    521         return _return_parsed_timezone_results(result, timezones, utc,
    ↪ name)

```

File `strptime.pyx:534`, in `pandas._libs.tslibs.strptime.array_strptime()`

File `strptime.pyx:355`, in `pandas._libs.tslibs.strptime.array_strptime()`

**ValueError:** time data "13/07/2013" doesn't match format "%m/%d/%Y", at position

↪ 9. You might want to try:

- passing ``format`` if your strings have a consistent format;
- passing ``format='ISO8601'`` if your strings are all ISO8601 but not
- ↪ necessarily in exactly the same format;
- passing ``format='mixed'``, and the format will be inferred for each element
- ↪ individually. You might want to use ``dayfirst`` alongside this.

#### Further Data Preparation for Modelling

```

[ ]: # List of categorical columns to one-hot encode
categorical_columns = ['City_Name', 'City_State', 'Climate', 'Season',
    ↪ 'Significant_Rainfall',
    ↪ 'Wind_Direction_AM', 'Wind_Direction_PM',
    ↪ 'Max_WindGust_Direction']

# Apply one-hot encoding to categorical columns
df_encoded = pd.get_dummies(df, columns=categorical_columns, drop_first=True)

# Print the first few rows of the encoded DataFrame
print(df_encoded.head())

```

```

[ ]: # List of categorical columns to one-hot encode
categorical_columns = ['City_Name', 'City_State', 'Climate', 'Season',
    ↪ 'Significant_Rainfall',
    ↪ 'Wind_Direction_AM', 'Wind_Direction_PM',
    ↪ 'Max_WindGust_Direction']

# Apply one-hot encoding to categorical columns
df_encoded = pd.get_dummies(df, columns=categorical_columns, drop_first=True)

# Print the first few rows of the encoded DataFrame
print(df_encoded.head())

```

```

[ ]: # Separate features and target variable
X = df_encoded.drop('Burn_Tomorrow', axis=1)
y = df_encoded['Burn_Tomorrow']

```



## Class Imbalance Handling

```
[ ]: # Step 1: Feature Scaling
scaler = StandardScaler()
# Exclude the 'Date' column from X
X_scaled = scaler.fit_transform(X.drop('Date', axis=1))

# Import the SMOTE class
from imblearn.over_sampling import SMOTE

# Instantiate SMOTE
smote = SMOTE(random_state=42)

# Apply SMOTE to the features and target variables
X_resampled, y_resampled = smote.fit_resample(X_scaled, y)

# Print the shape of the resampled data
print("Shape of X_resampled:", X_resampled.shape)
print("Shape of y_resampled:", y_resampled.shape)

# Split the resampled data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
    ↪test_size=0.2, random_state=42)

# Print the shape of train and test sets
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)
```

## Feature Selection

```
[ ]: from sklearn.ensemble import RandomForestClassifier

# Assuming X_train and y_train are your training data
model = RandomForestClassifier()
model.fit(X_train, y_train)

# Get feature importances
feature_importances = model.feature_importances_

# Rank features by importance
feature_ranking = sorted(range(len(feature_importances)), key=lambda i:
    ↪feature_importances[i], reverse=True)

# Select top N features (e.g., top 10)
top_n_features = feature_ranking[:10]
```

```

# Create a new dataset with the selected features
X_train_selected = X_train[:, top_n_features]
X_test_selected = X_test[:, top_n_features]

# Print the shape of train_selected and test_selected sets
print("Shape of X_train_selected:", X_train_selected.shape)
print("Shape of X_test_selected:", X_test_selected.shape)

```

```

[ ]: # List of feature names in the original dataset
feature_names = X.columns.tolist()

# Get the names of the top selected features
selected_feature_names = [feature_names[i] for i in top_n_features]

# Print the names of the selected features
print('selected_feature_names:\n', selected_feature_names)

```

## Training and Evaluation

```

[ ]: from sklearn.metrics import classification_report, accuracy_score

# Create a Random Forest Classifier
random_forest_classifier = RandomForestClassifier(random_state=42)

# Train the Random Forest Classifier on the selected features
random_forest_classifier.fit(X_train_selected, y_train)

# Perform cross-validation on the Decision Tree Classifier
scores = cross_val_score(random_forest_classifier, X_train_selected, y_train,
    ↪cv=5)

# Print the cross-validation scores
print("Cross-validation scores:", scores)

# Calculate the mean and standard deviation of the cross-validation scores
mean_score = scores.mean()
std_score = scores.std()

# Print the mean accuracy and standard deviation
print("Mean accuracy:", mean_score)
print("Standard deviation:", std_score)

# Make Predictions
y_pred_rf = random_forest_classifier.predict(X_test_selected)

# Evaluate the Random Forest Classifier
accuracy_rf = accuracy_score(y_test, y_pred_rf)

```

```
classification_report_rf = classification_report(y_test, y_pred_rf)

# Print the accuracy and classification report for Random Forest
print("Random Forest Classifier:")
print("Accuracy:", accuracy_rf)
print("Classification Report:")
print(classification_report_rf)
```

```
[ ]: confusion_matrix_rf = confusion_matrix(y_test, y_pred_rf)
print("Confusion Matrix:")
print(confusion_matrix_rf)
```

Copyright © 2025 Ningquan Du. All Rights Reserved.

This document is an academic assignment submitted to The University of Queensland.  
Any unauthorized use, reproduction, or distribution of this content is strictly prohibited.  
Plagiarism or any violation of academic integrity will be considered a serious offense.